

instruction pipelining and pre-computation in neural networks

colin shaw
09.08.2016

1. introduction

- a. welcome to computer science club
- b. thank you to revunit for sponsoring
- c. what we talked about the last two times
 - i. multiple layer perceptron with nice dsl in ocaml
 - ii. improving matrix arithmetic computational performance
 - iii. difference between functional and imperative programming
- d. what we are talking about this time
 - i. computer organization and speed optimization
 - 1. memory organization
 - 2. caching methods
 - 3. processor pipelining
 - 4. what environment makes numeric code fast
 - 5. code design to take advantage of cache and pipelining (including ocaml)
 - 6. wire, silicon and the speed of light
 - ii. neural network problem space reduction
 - 1. problem size vs training time in neural networks
 - 2. reducing feature space as pre-computation in neural networks
- e. why are we talking about this?
 - i. need to know how large a problem can be solved in known time
 - ii. need to know how large a problem training a mlp is
- f. why does that matter?
 - i. real time training
 - ii. low cost
 - iii. low power

2. memory organization

- a. alu / fpu
- b. processor registers (typically 8 - 16)
- c. I1i and I1d cache
 - i. instruction and data separate
 - ii. implemented per core
 - iii. fabricated on processor die
 - iv. sram
 - v. 0 - 64kb, power of two
 - vi. ≈ 1 ns access time
- d. I2 cache
 - i. blend of instruction and data
 - ii. higher end fabricated per core
 - iii. lower end shared between cores
 - iv. fabricated on processor die
 - v. sram
 - vi. 256kb - 1s of mb, power of two
 - vii. ≈ 5 ns access time
- e. I3 cache
 - i. blend of instruction and data
 - ii. shared between cores
 - iii. higher end fabricated on processor die (sram)
 - iv. lower end fabricated off processor die (dram)
 - v. 1s of mb - 10s of mb, generally power of two
 - vi. ≈ 30 ns access time
- f. I4 cache
 - i. mostly commonly implemented as edram for video
 - ii. focus is bandwidth rather than access time
 - iii. Embedded off processor die (edram)
 - iv. 10s of mb - 100s of mb, power of two
- g. main memory
 - i. blend of instruction and data
 - ii. shared
 - iii. off processor die (dram)
 - iv. 1s of gbs - 10s of gb
 - v. ≈ 100 ns access time
 - vi. 10,000+ mb/s (pc1333)
- h. solid state disk
 - i. 10,000 ns access time
 - ii. 500 mb/s
- i. rotational hard disk
 - i. 5,000,000 ns access time
 - ii. 200 mb/s

3. cache operation

- a. cache lines
 - i. basic unit of transport from cache to processor
 - ii. loading one byte in a line loads the whole line
 - iii. typically 16 bytes - 128 bytes (nominally 64)
- b. cache types
 - i. direct mapped
 - 1. each main memory location can go to exactly one cache location
 - 2. suffers contention in worst case
 - ii. fully associative
 - 1. each main memory location can map to any cache location
 - 2. costly to implement
 - iii. n-way associative
 - 1. compromise between direct mapped and fully associative
 - 2. each main memory location can go to n potential cache locations
 - 3. most common implementation of caching
 - 4. generally 2-way to 16-way associativity
 - 5. difference between intel and amd
- c. translation lookaside buffer (tlb)
 - i. address virtualization
 - ii. used as fast memory reference without computing physical page

4. processor pipelining

- a. key steps
 - i. read
 - ii. load
 - iii. execute
 - iv. write
- b. queueing
 - i. grocery checkout analogy
- c. how to control
 - i. modern processors have 12 - 24 stage pipeline
 - ii. branch and memory access prediction
 - iii. some assembly / compiler control
- d. multimedia and math extensions (intel)
 - i. mmx
 - ii. sse
 - iii. avx (256 bit registers)

5. what environment makes numeric code fast

- a. fast memory
 - i. dominated by main memory bandwidth
 - ii. sequential access best
 - iii. cache line alignment beneficial, especially for short sequential access
 - iv. buy the fastest you can afford
- b. registers and cache
 - i. enough registers for loops and intermediate values
 - ii. cache alignment
 - iii. prefetching and other memory optimizations
 - iv. buy the largest cache you can afford
 - v. pay attention to cache associativity
- c. fast processor
 - i. relative core speed to memory bandwidth
 - ii. generally involves pipelining and memory fetch consideration
 - iii. example raspberry pi vs intel processor
 - iv. buy the fastest core you can afford with the right pipelining
- d. intelligent choice of types
 - i. traditional numeric options
 - 1. single precision (32 bit)
 - 2. double precision (64 bit)
 - ii. more exotic numeric options
 - 1. long double (80 or 128 bit)
 - 2. half precision (16 bit)
 - a. what is it
 - i. generally a storage format with single computation
 - ii. better memory bandwidth
 - iii. small cost for conversion to/from single
 - iv. worse performance if sequence < l1d cache size
 - b. where is it seen in the wild
 - i. arm
 - ii. gpu
 - iii. intel generation 3 and above
 - iii. which is best and why
 - 1. smallest sufficient for problem
 - 2. best memory throughput
 - 3. best pipelining
- e. algorithm design
 - i. column-major vs row-major organization
 - ii. loop unrolling (often done better by compiler)
- f. compiler considerations
 - i. quality of loop unrolling
 - ii. quality of memory and cache consideration (prefetch, etc.)
 - iii. supported vectorization
 - iv. function inlining
 - v. type support (half precision, etc.)

6. coding examples for cache and pipelining

a. experiments

- i. `cache_line.c`
 - 1. cache line utilization vs aggregate speed
 - 2. x86 cache line is 64 bytes
 - 3. non-optimized and optimized (e.g. -O3)
 - 4. optimization is cache line alignment
- ii. `cache_speed.c`
 - 1. normalized time for sequential access by size
 - 2. detects L1 and L2 size
- iii. `pipeline.c`
 - 1. memory lookup pipelining and compiler optimization
 - 2. non-optimized and optimized
- iv. `pipeline2.c`
 - 1. instruction level pipelining
 - 2. non-optimized and optimized
 - 3. compiler may optimize two adds as one in tight loops
- v. `thread.c`
 - 1. pathological cache line access patterns with threading
 - 2. false cache line sharing

b. best practice patterns

- i. smallest type meeting problem criteria
- ii. tight loops
- iii. sequential accesses
- iv. function inlining
- v. compiler hinting
- vi. compiler optimization settings
- vii. compiler choice
 - 1. manufacturer with hand tuning
 - 2. general with optimizations
 - 3. general without optimizations

c. practical implementation with ocaml

- i. `ocaml_native.ml`
 - 1. intrinsic array map
- ii. `ocaml_no_opt.ml.ml / ocaml_no_opt.c.c`
 - 1. array map done in c
- iii. `ocaml_opt.ml.ml / ocaml_opt.c.c`
 - 1. bigarray
 - 2. in-place
 - 3. extra compiler optimizations for single precision floating point

7. wire, silicon and the speed of light

- a. speed of information in wire
 - i. 1 ns (1 ghz) \approx 1 foot propagation
 - ii. limitations
 - 1. distance from l1 cache to core at core speed
 - 2. distance from processor to dram distance at memory speed
- b. cmos process
 - i. mosfet gate area shrinks quadratically as lithography shrinks
 - ii. gate thickness single atom
 - 1. new gate material is new trend
 - 2. was silicon dioxide
 - 3. new materials harder harder to fab, more expensive
 - iii. gate capacitance becomes limiting
 - 1. rc circuit of gate becomes new prevailing issue
 - 2. vertical wire paths new technique
 - a. can make direct, euclidean distance, connections
 - b. most silicon connections are manhattan distance
 - 3. gates as amplifiers in path to circumvent rc time
 - iv. heat generation
 - 1. moving charges generate heat
 - 2. faster clock yields more heat
 - 3. higher density yields higher temperature rise
 - 4. must be at least somewhat above bandgap to work (limit)
- c. notes on lithography
 - i. current production technology
 - 1. 14nm
 - a. smaller than smallest known virus
 - b. 40 times shorter than wavelength of green light
 - 2. 5nm by early 2020s
 - ii. how it works
 - 1. (extreme) ultraviolet photolithography
 - 2. wavelength (slightly) smaller than features
 - 3. diffraction effects specifically used
 - 4. surface roughness is problematic (atomic surface is 0.23 nm)
 - 5. mask and etch, mask and etch...

8. neural network problem size and training time

- a. theory
 - i. problem degrees of freedom
 - 1. total inputs and outputs
 - 2. depends on this being proper size for problem
 - a. cannot be underdetermined or will have training problems
 - b. cannot be overdetermined or will not be good estimator
 - ii. hidden layer effects
 - 1. each neuron is independent per layer
 - 2. each layer is independent because of nonlinear activation function
 - 3. hidden layer neurons are multipliers of problem dof
 - iii. putting it all together
 - 1. $n_t = \alpha (n_i + n_o) n_h$
- b. simplification
 - i. assumptions
 - 1. $n_i = v^2$
 - 2. $n_o = k_1$
 - 3. $n_h = k_2 v$
 - ii. estimation
 - 1. $n_t = \alpha (v^2 + k_1) k_2 v$
 - 2. n_t bounded by $\alpha k_2 v^3$
 - iii. example
 - 1. assumptions
 - a. $n_t = 2000$
 - b. $\alpha = 2$
 - c. $k_2 = 1$
 - 2. $v \cong 10$

9. pre-processing in neural networks

- a. rationale
 - i. biological imperative - retina
 - ii. problem determinism
 - iii. improved performance
 - iv. improved understanding of problem
- b. an example for detecting sidewalk curvature
 - i. create subsampled images from original feature set
 - ii. threshold images to black and white
 - iii. image output is zero if not enough blend of black and white
 - iv. find centroid of black and white
 - v. image vector between centroids
 - vi. create perpendicular vector, indicating tangent of curve
 - vii. image output is angle of tangent
 - viii. expectation is zeros except sides of sidewalk